

# Introduction to Applied Bayesian Statistics

## Part III: Building Bayes Theorem

## Part IV: Prior Specification

Michael Anderson, PhD

Department of Biostatistics and Epidemiology  
The University of Oklahoma Health Sciences Center

Sept 14, 2016

# Outline

## Example

In 2012, the prevalence of HIV among the general U.S. population was estimated to be 0.38% ([www.cdc.gov](http://www.cdc.gov)). Rapid diagnostic tests have been developed to test for the presence of HIV infection in as little as 10 minutes. Suppose one such diagnostic test was evaluated in a case-control study and it was observed that out of 1000 subjects with HIV, 967 obtained positive test results and out of 1000 subjects without HIV, 985 obtained negative test results.


### In summary

- $P(D+) = 0.0038 \Rightarrow P(D-) = 1 - 0.0038 = 0.9962$
- $P(T+ | D+) = 0.967 \Rightarrow P(T- | D+) = 1 - 0.967 = 0.033$
- $P(T- | D-) = 0.985 \Rightarrow P(T+ | D-) = 1 - 0.985 = 0.015$


Suppose a person in the U.S. is selected at random, given the test, and the result is positive. Should this person begin treatment?

- $P(D+ | T+)$ ?

# Venn Diagram of the Example



# Venn Diagram of the Example



Note

- $P(D+ \cap T+) = P(D+)P(T+ | D+)$
- $P(D- \cap T+) = P(D-)P(T+ | D-)$
- $P(T+) = P(D+)P(T+ | D+) + P(D-)P(T+ | D-)$

$$P(D+ | T+) = \frac{P(D+ \cap T+)}{P(T+)} = \frac{P(D+)P(T+ | D+)}{P(D+)P(T+ | D+) + P(D-)P(T+ | D-)}.$$

## Example (continued)


Recall

- $P(D+) = 0.0038 \Rightarrow P(D-) = 1 - 0.0038 = 0.9962$
- $P(T+|D+) = 0.967 \Rightarrow P(T-|D+) = 1 - 0.967 = 0.033$
- $P(T-|D-) = 0.985 \Rightarrow P(T+|D-) = 1 - 0.985 = 0.015$


$$P(D+|T+) = \frac{P(D+)P(T+|D+)}{P(D+)P(T+|D+)+P(D-)P(T+|D-)} =$$
$$\frac{(0.0038)(0.967)}{(0.0038)(0.967)+(0.9962)(0.015)} = 0.1974$$

What does this mean?

# More Generally...




# More Generally...







# More Generally...




# More Generally...




# More Generally...




# More Generally...




# More Generally...




# More Generally...



# More Generally...



# More Generally...





## More Formally...

### Bayes' Theorem

Let  $A_1, \dots, A_k$  be a partition of the sample space, and let  $B$  be any event in  $S$ . Then, for each  $i = 1, \dots, k$

$$P(A_i|B) = \frac{P(A_i)P(B|A_i)}{\sum_{j=1}^k P(A_j)P(B|A_j)} \quad (1)$$

- **Prior** probability of  $A_i$ .
- **Conditional** of  $B$  given  $A_i$ .
- **Total Probability**.
- **Posterior** probability of  $A_i$ .

# Joint Probability Function

Suppose a *sample* is made up of independent observations  $X_1, \dots, X_n$  all assumed to belong to the same identical pdf (or pmf)  $p(x|\theta)$ . Then  $X_1, \dots, X_n$  are said to be i.i.d. (independent and identically distributed).

## Joint pdf (or pmf)

The joint pdf (or pmf) of an i.i.d. *sample*  $X = (X_1, \dots, X_n)$  is given by

$$p(x_1, \dots, x_n | \theta) = \prod p(x_i | \theta)$$

# Likelihood

## Likelihood Function

Let  $p(x_1, \dots, x_n | \theta)$  denote the joint pdf or pmf of the *sample*  $X = (X_1, \dots, X_n)$ . Then, given  $x_1, \dots, x_n$  is observed, the function of  $\theta$  defined by

$$L(\theta | x_1, \dots, x_n) = p(x_1, \dots, x_n | \theta)$$

is called the *likelihood function*

Plotting this function vs  $\theta$  shows how plausible each  $\theta$  value is.

The maximum of the likelihood function is seen as the most plausible value of  $\theta$ , given the data that was observed.

## More Generally...

### Bayes' Theorem

$$p(\theta|x_1, \dots, x_n) = \frac{p(\theta)p(x_1, \dots, x_n|\theta)}{\int p(\theta)p(x_1, \dots, x_n|\theta)d\theta} \quad (2)$$

- **Prior** distribution of  $\theta$ .
- **Likelihood** of the data.
- Normalizing **constant**.
- **Posterior** distribution of  $\theta$ .
- Not always tractable.
- Conjugate Priors.
- Gibbs Sampling and the Metropolis Algorithm.

# Mean (Expected Value) of a Probability Distribution

- Discrete:  $E(x) = \sum_x p(x)x$ 
  - Roll a single fair die infinitely many times. What is the mean of all the rolls?

$$E(x) = (1/6)1 + (1/6)2 + (1/6)3 + (1/6)4 + (1/6)5 + (1/6)6 = 3.5$$

- Continuous:  $E(x) = \int_x p(x)x dx$ 
  - For a Normally distributed (continuous) outcome

$$E(x) = \int_x \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) x dx = \mu$$


# Variance of a Probability Distribution

- Discrete:  $Var(x) = \sum_x p(x)(x - E(x))^2$
- Continuous:  $Var(x) = \int_x p(x)(x - E(x))^2 dx$
- Note this is just the expected value of  $(x - E(x))^2$ .
- The positive square root of the variance is the standard deviation.
- **Major point:** The variance can represent our uncertainty about possible beliefs.
  - Less uncertainty about beliefs  $\Rightarrow$  smaller variance.
  - Less certain about beliefs  $\Rightarrow$  larger variance.

# Bernoulli Distribution

- A Bernoulli experiment consists of a single trial with two possible outcomes (success/failure) with success probability  $\pi$ .
- $Bern(\pi)$
- $p(x) = \pi^x(1 - \pi)^{1-x} \quad x = 0, 1$
- $E(x) = \pi$
- $Var(x) = \pi(1 - \pi)$
- Example: Plot a  $Bern(0.3)$ 
  - `x<-c(0,1)`
  - `plot(x,dbinom(x,1,0.3),ylim=range(0,1),type="h",  
ylab="Probability",xlab="Success/Failure",col="blue")`

# Bernoulli Distribution






# Binomial Distribution

- A Bernoulli experiment repeated  $n$  times.
- Outcome of interest is the number of successes in those trials.
- $\text{Bin}(n, \pi)$
- $p(x) = \binom{n}{x} p^x (1 - p)^{n-x} \quad x = 0, 1, \dots, n$
- $E(x) = n\pi$
- $\text{Var}(x) = n\pi(1 - \pi)$
- Example: Plot a  $\text{Bin}(25, .3)$ 
  - `x<-0:25`
  - `plot(x,dbinom(x,25,.3),ylim=range(0,.25),type="h",  
ylab="Probability",xlab="Number of  
Successes",col="blue")`


# Binomial Distribution



# Discrete Uniform Distribution

- Puts equal probability on every observable outcome.
- $DUniform(N)$
- $p(x) = \frac{1}{N} \quad x = 1, 2, \dots, N$
- $E(x) = \frac{N+1}{2}$
- $Var(x) = \frac{(N+1)(N-1)}{12}$
- Example: Plot a discrete uniform for the roll of a fair die.
  - `x<-1:6`
  - `plot(x,rep(1/length(x),length(x)),ylim=range(0,1),type="h",  
ylab="Probability",xlab="Possible Outcome  
Values",col="blue")`


# Discrete Uniform Distribution



# Uniform Distribution

- Puts equal density on every subinterval of the same length between two points  $[a, b]$ .
- $Uniform(a, b)$
- $p(x) = \frac{1}{b-a} \quad a \leq x \leq b$
- $E(x) = \frac{b+a}{2}$
- $Var(x) = \frac{(b-a)^2}{12}$
- Example: Plot a Uniform(0,1)
  - `x<-seq(0,1,length=1000)`
  - `plot(x,dunif(x,0,1),ylim=range(0,1),type="l",  
ylab="Density",xlab="Possible Outcome  
Values",col="blue")`


# Uniform Distribution



# Normal Distribution

- Symmetric, bell-shaped curve.
- Outcome of interest could have any value in the real numbers.
- $N(\mu, \sigma^2)$
- $p(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) \quad -\infty < x < \infty$
- $E(x) = \mu$
- $Var(x) = \sigma^2$
- Example: Plot a  $N(0,1)$ 
  - `x<-seq(-4,4,length=1000)`
  - `plot(x,dnorm(x,0,1),ylim=range(0,.25),type="l",  
ylab="Density",xlab="Possible Outcome  
Values",col="blue")`

# Normal Distribution






# Gamma Distribution

- Sometimes used to model lifetimes. Usually right-skewed. Often used as a prior for  $1/\text{Var}(x)$ .
- Outcome of interest must be non negative.
- $\text{Gamma}(\alpha, \beta)$
- $p(x) = \frac{1}{\Gamma(\alpha)\beta^\alpha} x^{\alpha-1} e^{-\frac{x}{\beta}} \quad 0 < x < \infty$
- $E(x) = \alpha\beta$
- $\text{Var}(x) = \alpha\beta^2$
- Example: Plot a  $\text{Gamma}(5,2)$ 
  - `x<-seq(0,40,length=1000)`
  - `plot(x,dgamma(x,5,.5),ylim=range(0,.15),type="l",  
ylab="Density",xlab="Possible Outcome  
Values",col="blue")`

# Gamma Distribution



# Beta Distribution

- Often used to model probabilities or prevalences.  
Often used as a prior for these quantities.

- Outcome of interest lives on the interval  $[0, 1]$ .

- $Beta(\alpha, \beta)$

- $p(x) = \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1} \quad 0 < x < \infty$


- $E(x) = \frac{\alpha}{\alpha+\beta}$

- $Var(x) = \frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$


- Example: Plot a Beta(3,5)

- `x<-seq(0,1,length=1000)`
- `plot(x,dbeta(x,3,5),ylim=range(0,3),type="l",  
ylab="Density",xlab="Possible Outcome  
Values",col="blue")`

# Beta Distribution



# Beta Distribution



# Conjugate Priors

Posterior has the same distributional form as the prior, then the prior is conjugate for the likelihood.

**Table:** Common Conjugate Priors and Corresponding Likelihoods.

Likelihood	Prior	Posterior
Bernoulli	Beta	Beta
Binomial	Beta	Beta
Normal	Normal	Normal
Poisson	Gamma	Gamma

# Refueling Example Revisited

Table: Results of 24 refueling attempts;  
1=Premature Disconnect

Status	Status
0	0
1	0
0	0
0	1
1	1
0	0
0	0
0	1
0	0
0	1
1	0
0	0

- Note that  $x_i \in \{0, 1\}$
- Each observation  $x_i \sim \text{Bern}(\theta)$
- $p(x_i|\theta) = \theta^{x_i}(1-\theta)^{1-x_i}$
- $L(\theta|x_i) = \theta^{\sum x_i}(1-\theta)^{n-\sum x_i}$
- To ease notation let  $y = \sum x_i$

- Prior:  $p(\theta) = \text{Beta}(\alpha, \beta) = \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)}\theta^{\alpha-1}(1-\theta)^{\beta-1}$
- Mean:  $\frac{\alpha}{\alpha+\beta}$
- Variance:  $\frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$







# Refueling Example Revisited

- Posterior:  $P(\theta|y) = \frac{p(\theta)p(y|\theta)}{\int p(\theta)p(y|\theta)d\theta} = \frac{\theta^{\alpha-1}(1-\theta)^{\beta-1}\theta^y(1-\theta)^{n-y}}{\int \theta^{\alpha-1}(1-\theta)^{\beta-1}\theta^y(1-\theta)^{n-y}d\theta}$
- $= \frac{\theta^{\alpha-1+y}(1-\theta)^{\beta-1+n-y}}{\int \theta^{\alpha-1+y}(1-\theta)^{\beta-1+n-y}d\theta} = \frac{\Gamma(y+\alpha+n-y+\beta)}{\Gamma(y+\alpha)\Gamma(n-y+\beta)} \theta^{y+\alpha-1}(1-\theta)^{n-y+\beta-1}$
- $= \text{Beta}(\alpha + y, n - y + \beta)$
- Mean:  $\frac{\alpha+y}{\alpha+\beta+n} = \left(\frac{\alpha+\beta}{\alpha+\beta+n}\right) \left(\frac{\alpha}{\alpha+\beta}\right) + \left(1 - \frac{\alpha+\beta}{\alpha+\beta+n}\right) \left(\frac{y}{n}\right)$
- Variance:  $\frac{(\alpha+y)(n-y+\beta)}{(\alpha+\beta+n)^2(\alpha+\beta+2n)}$



# Refueling Example Revisited

Consider a  $Beta(1, 1)$  and  $Beta(15, 15)$  as possible priors



- Mean : 0.308
- Var : 0.008

0.407  
0.004

# Prior Beliefs

- Knowledge about parameters (or common data biases) BEFORE any data is observed.

We rarely conduct experiments without some degree of prior belief.

- Hypotheses of statistical tests come from prior beliefs.
- Sample size calculations require prior information.
- Identification of outliers depends on what was expected *a priori*.
- Anticipated bias in the data could be handled with prior knowledge of the bias.

# Eliciting Prior Beliefs

The goal of eliciting prior beliefs is to turn them into a probability function for use in analysis.

## Hierarchy of Prior Knowledge

- Meta analysis of the topic from literature.
- Previous literature of single study.
- Experts' opinion.-Visualizations such as the online MATCH tool are very useful.

In the absence of any of the above sources one may use diffuse (vague) or non-informative priors.

# Using the MATCH Uncertainty Elicitation Tool

## MATCH Uncertainty Elicitation Tool

- Step 1. From the homepage select the range of possible values using the Upper and Lower limits.
- Step 2. Select Roulette Input Mode. There are other input modes but this one is easy and clever.
- Step 3. Place chips in bins along the specified interval to reflect prior beliefs.
- Step 4. Click “Fitting & Feedback” to see best fitting distributions with accompanying parameter values.

# Using the MATCH Uncertainty Elicitation Tool

